Aalto University
School of Science
Degree programme in Engineering Physics and Mathematics

Comparison of line search methods in unconstrained optimization

Bachelor's Thesis
1.5.2019

Einari Tuukkanen

AALTO-YLIOPISTO PERUSTIETEIDEN KORKEAKOULU PL 11000, 00076 Aalto http://www.aalto.fi	KANDIDAATINTYÖN TIIVISTELMÄ	

Tekijä: Teppo Teekkari

Työn nimi: Teekkarien valmistumistodennäköisyyden mallintaminen Markov-ketjuilla

Tutkinto-ohjelma: Teknillisen fysiikan ja matematiikan tutkinto-ohjelma

Pääaine: Systeemitieteet	Pääaineen koodi: F3010

Vastuuopettaja(t): Prof. Ansio Akateemikko

Ohjaaja(t): DI Otto Operaatioinsinööri

Tiivistelmä:

Tämä on LaTeX-pohjaan kuuluva esimerkki tiivistelmäsivusta. Korvaa tämä tiedosto omalla tiivistelmälläsi, jonka olet tallentanut esim. Microsoft Officesta tai LibreOfficesta PDF-muotoon. Uusimman tiivistelmäpohjan löydät TFM-kandidaatintyön ja seminaarin Noppa-sivuilta.

Päivämäärä: 11.11.2011	Kieli: suomi	Sivumäärä: 4+1

Avainsanat: esimerkki, lorem ipsum

Contents

1 Introduction

There are a number of sources for real-life unconstrained nonlinear optimization problems from different fields of sciences and technologies. In general an unconstrained optimization problem can be formulated as

$$\min_{\mathbf{x} \in \mathbb{R}^n} f(\mathbf{x}). \tag{1}$$

Depending on the methods used, there may also be additional requirements for f, such as being continuously differentiable up to n times or being convex i.e. having one unique global minimum. While there are a number of real-life problems that do satisfy these conditions they are often extremely complex. Therefore artificial problems are often used to develop and test optimization algorithms. [1]

The optimization algorithms are iterative processes that abuse the properties of the target function combined with the raw processing powers of computers. The commonly used methods are based on the same principle of having some termination condition checking if the minimum was found, and if not, continuing by choosing a direction and a step size and moving to the next point and repeating. The step size is often chosen by a separate algorithm called a line search, which transforms the target function into one-dimensional function of step size λ and using the new function to calculate the optimal step size for each iteration of the main method. Similarly to the main methods, there are a number of different algorithms for finding the optimal step size, the simplest being just using a constant λ for every iteration.

In this thesis we will be exploring the effects of line search method selection on the overall performance of the optimization method. The thesis focues primarily on the use cases introduced on the Nonlinear Optimization (2018-2019) course in the Aalto University. Since this bachelor's thesis will not be completely comprehensive, the main goal will be answering the question, will the selection of line searches effect the main method's performance and if so, by how much? In addition we will be making a comparison of performance differences caused by the line searches in a very specific testing setup with goal of providing a suggestion of how the line search methods could be ranked.

> Should I mention the target functions, main methods and line search methods explored already here on intro?

> What is the recommended style for e.g. marking vectors?
> Which equations should I numerate?

Abbreviations and symbols used

Table 1: Abbreviations used in this thesis, and their explanations.

Abbreviation	Explanation
MM	main method, an unconstrained optimization method
LS	line search, provides step sizes for main methods
NM	Newton's Method
GDM	Gradient Descent Method
CGM	Conjugate Gradient Method
HBM	Heavy Ball Method
CS	Constant search i.e. constant step size
GSS	Golden Section line search
BS	Bisection line search
DS	Dichotomous line search
FS	Fibonacci line search
US	Uniform line search
NS	Newton's line search
AS	Armijo's line search
MSS	Matrix Square Sum target function
NE	Negative Entropy target function

2 Theory

2.1 Target Functions

2.1.1 Choosing target functions

One of the biggest factors effecting the performance of an optimization algorithm is the complexity of the target function we are using the method for. In this thesis we are going to focus on two different functions. Both of the functions are convex and differentiable since some of the methods we will be comparing require these

properties. To increase the complexity of the problems a bit, we are going to use 50-dimensional versions of the two functions. Therefore we are going to pick functions that are defined in $\mathbb{R}^n$ or similar domain and then set $n = 50$.

2.1.2 Step size function

Each optimization method involves a step in form of $\operatorname{argmin}_{\lambda \in \mathbb{R}} f(x_i + \lambda d_i)$. To programmatically produce this step, we are going to define step size function $g(\lambda)$ related to each of the target functions. Then on each iteration of the main optimization method, a new step size function is seeded with the updated point $\mathbf{x}$ and direction vector $\mathbf{d}$.

2.1.3 Matrix Square Sum

The Matrix Square Sum is an extended version of the regular Sum Squares Function, which is defined as

$$f(\mathbf{x}) = f(x_1, ..., x_n) = \sum_{i=1}^{n} i x_i^2. \tag{2}$$

To make Sum Squares problem less trivial, we are going modify it by adding three constants A, b and c. Together they form what we are going to call a Matrix Square Sum function, defined as

$$f(\mathbf{x}) = \|A\mathbf{x} + \mathbf{b}\|^2 + c\|\mathbf{x}\|^2, \tag{3}$$

where A is a positive definite (PD) $n \times n$ matrix, $\mathbf{x}$ and $\mathbf{b}$ are vectors of length n, and the scalar c is a positive constant.

While A could in theory be any $m \times n$ matrix, we are limiting it to a square PD-matrices only to ensure that the function is convex and has an unique global minimum. The matrix A is formed in few steps by first generating an initial $n \times n$ matrix A' with random values $a_{ij} \in [-0.5, 0.5] \,\forall\, i, j = 1 \ldots n$. Then we attempt to make a PD-matrix by $A = 0.5(A' + A'^\top)$. Finally, if A does not happen to be a PD-matrix yet, it is modified by the formula

$$A = A + (|\lambda_{min}| + d)I, \tag{4}$$

where λ_{min} is the minimal eigenvalue of A and d is a constant. The value of d also determines the A's condition number defined as $\frac{\lambda_{max}}{\lambda_{min}}$. Smaller values of d make the function more elliptic and harder for gradient methods to solve, while larger values produce more circular gradient curves, which are easy for methods

like Gradient Descent to solve. The value of d used in the thesis is 5, resulting in condition number of around 6.5.

The values of b and c are also randomized so that for each element of b applies $b_i \in [-0.5, 0.5]$ and c is a scalar randomly generated from range $[-0.5, 0.5]$.

Some of the optimization methods require access to the function's gradient, Hessian matrix and the step size function with its first two derivatives, so let's calculate them in advance. While at it, we should also derive the formula for calculating the correct optima so we can use it to determine whether our algorithms produce the correct solution. Lets begin by expanding the function into more easily differentiable form:

$$
\begin{aligned}
f(\mathbf{x}) &= \|A\mathbf{x} + \mathbf{b}\|^2 + c\|\mathbf{x}\|^2 \\
&= (A\mathbf{x} + \mathbf{b})^\top (A\mathbf{x} + \mathbf{b}) + c\mathbf{x}^\top \mathbf{x} \\
&= (\mathbf{x}^\top A^\top + \mathbf{b}^\top)(A\mathbf{x} + \mathbf{b}) + c\mathbf{x}^\top \mathbf{x} \\
&= \mathbf{x}^\top A^\top A\mathbf{x} + \mathbf{x}A^\top \mathbf{b} + \mathbf{b}^\top A\mathbf{x} + \mathbf{b}^\top \mathbf{b} + c\mathbf{x}^\top \mathbf{x} \\
&= \mathbf{x}^\top A^\top A\mathbf{x} + 2\mathbf{b}^\top A\mathbf{x} + \mathbf{b}^\top \mathbf{b} + c\mathbf{x}^\top \mathbf{x}.
\end{aligned}
$$

Now let's calculate the gradient and then Hessian matrix:

$$
\begin{aligned}
\nabla f(\mathbf{x}) &= \nabla(\mathbf{x}^\top A^\top A\mathbf{x}) + 2\nabla(\mathbf{b}^\top A\mathbf{x}) + \nabla(\mathbf{b}^\top \mathbf{b}) + \nabla(c\mathbf{x}^\top \mathbf{x}) \\
&= (\mathbf{x}^\top A^\top A)^\top + A^\top A\mathbf{x} + 2(\mathbf{b}^\top A)^\top + c(\mathbf{x}^\top)^\top + c\mathbf{x} \\
&= 2A^\top A\mathbf{x} + 2A^\top \mathbf{b} + 2c\mathbf{x}
\end{aligned}
\tag{5}
$$

$$
\mathbf{H} = \nabla^2 f(\mathbf{x}) = 2A^\top A + 2cI
\tag{6}
$$

From the gradient we get the necessary optimality condition

$$
(A^\top A + cI)\mathbf{x} = -A^\top \mathbf{b}
\tag{7}
$$

and because the function is convex (the Hessian is positive definite for all $\mathbf{x} \in \mathbb{R}^n$), we get the unique optimal solution

$$
\mathbf{x} = -(A^\top A + cI)^{-1} A^\top \mathbf{b}
\tag{8}
$$

Let's now form the step size function and its derivatives:

$$
\begin{aligned}
g(\lambda) = f(\mathbf{x} + \lambda \mathbf{d}) &= \|A(\mathbf{x} + \lambda \mathbf{d}) + \mathbf{b}\|^2 + c\|(\mathbf{x} + \lambda \mathbf{d})\|^2 \\
&= \lambda \mathbf{d}^\top A^\top (2A\mathbf{x} + \lambda A\mathbf{d} + 2\mathbf{b}) + \mathbf{x}^\top A^\top (A\mathbf{x} + 2\mathbf{b}) + \mathbf{b}^\top \mathbf{b} \\
&\quad + c(\mathbf{x}^\top \mathbf{x} + 2\lambda \mathbf{d}^\top \mathbf{x} + \lambda^2 \mathbf{d}^\top \mathbf{d}) \\
g'(\lambda) &= \mathbf{d}^\top A^\top (2A\mathbf{x} + 2\lambda A\mathbf{d} + 2\mathbf{b}) + 2c\mathbf{d}^\top (\mathbf{x} + \lambda \mathbf{d}) \\
g''(\lambda) &= 2\mathbf{d}^\top (A^\top A\mathbf{d} + c\mathbf{d})
\end{aligned}
$$

$$(9)$$
$$(10)$$
$$(11)$$

2.1.4 Negative Entropy

The second target function reviewed is a variation of entropy maximization problem, in which the maximation problem is inverted therefore making it a minimization problem. We will be calling this convex minimization problem a Negative Entropy problem and it can be formulated as follows:

$$
f(\mathbf{x}) = \sum_{i=1}^n x_i \log x_i, \tag{12}
$$

where $\mathbf{x} \in \mathbb{R}^n_{>0}$. Since the domain is now limited we have to do some adjustments to our algorithms, which we will get into in a moment.

The global minima of the function can be found by calculating the zero point of the gradient:

$$
\nabla f(\mathbf{x}) = \frac{\delta}{\delta x_i} \sum_{i=1}^n x_i \log x_i \tag{13}
$$

$$
= \sum_{i=1}^n (\log x_i + 1) = 0. \tag{14}
$$

Because the $\log x_i + 1$ is truly positive, the answer is the same for every term x_i. Therefore we get the minima: $\log x_i = -1 \Leftrightarrow x_i = \frac{1}{e}$, for every $i = 1 \ldots n$.

In addition to the gradient, we need to calculate the Hessian matrix and the one dimensional step size function $g(\lambda) = f(\mathbf{x} + \lambda \mathbf{d})$ with its first and second derivatives in regard to λ.

As we already know the gradient of f, the Hessian can be derived with ease since $\frac{\delta^2}{\delta x_i \delta x_j} \sum_{i=1}^n x_i \log x_i = 0$ for every i, j where $i \neq j$. For the diagonal cases, i.e. when $i = j$, we are left with the second derivative $\frac{1}{x_i}$. Therefore we get the following

Hessian matrix:

$$\mathbf{H} = \begin{bmatrix} x_1^{-1} & 0 & 0 & \dots & 0 \\ 0 & x_2^{-1} & 0 & \dots & 0 \\ \vdots & \vdots & \vdots & \ddots & \vdots \\ 0 & 0 & 0 & \dots & x_n^{-1} \end{bmatrix}. \tag{15}$$

Finally, the step size function for λ with its first two derivatives are as follows:

$$g(\lambda) = \sum_{i=1}^{n}(x_i + \lambda d_i)\log(x_i + \lambda d_i) \tag{16}$$

$$g'(\lambda) = \sum_{i=1}^{n} d_i(\log(x_i + \lambda d_i) + 1) \tag{17}$$

$$g''(\lambda) = \sum_{i=1}^{n} \frac{d_i^2}{x_i + \lambda d_i}. \tag{18}$$

Because the domain of this function is only positive numbers, we are going to have make some changes to our starting points, parameters and algorithms. Therefore the starting point must only have positive coordinates and the following domain limiter was added to the line search algorithms which update the starting point so that it could otherwise go negative:

Algorithm 1 Domain Limiter

1: **initialize** $\gamma \in [0, 1]$
2: **while** min $(\mathbf{x} + \lambda\mathbf{d}) <= 0$ **do**
3:    $\lambda = \lambda\gamma$
4: **end while**
5: **return** λ

The domain limiter will be called from all of the line search methods before any evaluation of $f(\lambda)$ to reduce the step size to such value that it prevents evaualation of negative logarithms. In the context of this thesis we are going to set $d = 0.99$, but it could also be added as a parameter for each of the line search functions as some methods will require more iterations on the while-loop than others.

2.2 Main Methods

To make a difference between the unconstrained optimization methods and line search methods, we are going to call the first mentioned as main methods in this

thesis.

2.2.1 Newton's Method

Newton's method is commonly used for finding global minimums for differentiable and convex functions. The algorithm is a more general version of its line search version, utilizing the gradient and hessian matrix of the function instead of its first and second single derivatives.

Algorithm 2 Newton's Method

1: **initialize** $l > 0, k = 0$, $k_{max} \in \mathbb{N}$, $\mathbf{x} = \mathbf{x_0} \in \mathbb{R}^n$, step size function $g_{\mathbf{x},\mathbf{d}}(\lambda)$ and line search method L.
2: **while** $\|\nabla\theta(\mathbf{x})\| > l$ **and** $k < k_{max}$ **do**
3:     $\mathbf{d} = -H(\mathbf{x})^{-1}\nabla\theta(\mathbf{x})$
4:     $\lambda = L(g_{\mathbf{x},\mathbf{d}})$
5:     $\mathbf{x} = \mathbf{x} + \lambda\mathbf{d}$
6:     $k = k + 1$
7: **end while**
8: **return** λ

2.2.2 Gradient Descent Method

Gradient descent method can be interpret as a first order of the Newton's method as it only uses the first gradient to determine step direction. The algorithm itself is exactly the same as in Newton's method except for the line 3 where the descent direction d is calculated.

Algorithm 3 Gradient Descent Method

1: **initialize** $l > 0, k = 0$, $k_{max} \in \mathbb{N}$, $\mathbf{x} = \mathbf{x_0} \in \mathbb{R}^n$, step size function $g_{\mathbf{x},\mathbf{d}}(\lambda)$ and line search method L.
2: **while** $\|\nabla\theta(\mathbf{x})\| > l$ **and** $k < k_{max}$ **do**
3:     $\mathbf{d} = -\nabla\theta(\mathbf{x})$
4:     $\lambda = L(g_{\mathbf{x},\mathbf{d}})$
5:     $\mathbf{x} = \mathbf{x} + \lambda\mathbf{d}$
6:     $k = k + 1$
7: **end while**
8: **return** λ

2.2.3 Conjugate Gradient Method

Conjugate gradient method uses the information from the previous and current iterations to calculate ratio of gradients to calculate optimal next point. The method should theoretically require less steps than gradient descent method. source?

The algorithm involves a second loop which runs for n times per every iteration on the main loop. The n is equal to the dimensions of the problem (i.e. $\dim(\mathbf{x})$) and also the inner iterations are calculated towards the total iterations.

Algorithm 4 Conjugate Gradient Method

1: **initialize** $l > 0, k = 0$, $k_{max} \in \mathbb{N}$, $\alpha \in \mathbb{R}$, $\mathbf{x} = \mathbf{x_0} \in \mathbb{R}^n$, step size function $g_{\mathbf{x},\mathbf{d}}(\lambda)$ and line search method L.
2: $\mathbf{d} = -\nabla\theta(\mathbf{x})$
3: **while** $\|\nabla\theta(\mathbf{x})\| > l$ **and** $k < k_{max}$ **do**
4: $\quad$ $\mathbf{y} = \mathbf{x}$
5: $\quad$ **for** $i = 1 \ldots n$ **do**
6: $\quad\quad$ $\lambda = L(g_{\mathbf{y},\mathbf{d}})$
7: $\quad\quad$ $\mathbf{y_{prev}} = \mathbf{y}, \;\; \mathbf{y} = \mathbf{y} + \lambda\mathbf{d}$
8: $\quad\quad$ $\alpha = \frac{\nabla\theta(\mathbf{y})^2}{\nabla\theta(\mathbf{y_{prev}})^2}$
9: $\quad\quad$ $\mathbf{d} = -\nabla\theta(\mathbf{y}) + \alpha\mathbf{d}$
10: $\quad\quad$ $k = k + 1$
11: $\quad$ **end for**
12: $\quad$ $\mathbf{x} = \mathbf{y}, \;\; \mathbf{d} = \nabla\theta(\mathbf{x})$
13: **end while**
14: **return** λ

2.2.4 Heavy Ball Method

Heavy ball method is basically an extension of gradient descent as only adds one term for the direction vector calculation step. The name of the method comes from the physical representation of momentum of a ball rolling downhill. The algorithm source? for the method is basically the same as Gradient descent with a single additional term with multiplier β in descent direction d.

Algorithm 5 Heavy Ball Method

1: **initialize** $l > 0, k = 0$, $k_{max} \in \mathbb{N}$, $\beta \in \mathbb{R}$, $\mathbf{x} = \mathbf{x_{prev}} = \mathbf{x_0} \in \mathbb{R}^n$, step size function $g_{\mathbf{x},\mathbf{d}}(\lambda)$ and line search method L.
2: **while** $\|\nabla\theta(\mathbf{x})\| > l$ **and** $k < k_{max}$ **do**
3:     $\mathbf{d} = -\nabla\theta(\mathbf{x}) + \beta(\mathbf{x} - \mathbf{x_{prev}})$
4:     $\lambda = L(g_{\mathbf{x},\mathbf{d}})$
5:     $\mathbf{x_{prev}} = \mathbf{x}, \ \ \mathbf{x} = \mathbf{x} + \lambda\mathbf{d}$
6:     $k = k + 1$
7: **end while**
8: **return** λ

2.3   Line Search Methods

Most line search methods are based on a single idea, which we are going to call the theorem 1:

Theorem 1: Let $\theta : \mathrm{R} \rightarrow \mathrm{R}$ be strictly quasiconvex over the interval $[a, b]$, and let $\lambda, \mu \in [a, b]$ such that $\lambda < \mu$. If $\theta(\lambda) > \theta(\mu)$, then $\theta(z) \geq \theta(\mu)$ for all $z \in [a, \lambda]$. If $\theta(\lambda) \leq \theta(\mu)$, then $\theta(z) \leq \theta(\lambda)$ for all $z \in [\mu, b]$. [2]

2.3.1   Constant Step Size

While not actually a search, the use of single step size value λ is sometimes enough to provide correct and high performance solutions for the main methods. In this thesis I will be considering constant λ as one of the search methods to provide comprehensive comparison.

Table 2: Parameters tested for ConstantSearch

Parameter	MatrixSquareSum	NegativeEntropy
λ	0.0001, 0.1, 0.25, 0.5, 0.9	0.1, 0.25, 0.5

Because Constant Search has some convergence issues with some of the configurations with the test values listed in table 2, a different set of values must be used for some main methods. For Matrix Square Sum and Newton's Method the test values of parameter λ are $\{0.5, 0.9, 1.0, 1.1, 1.5\}$. For Negative Entropy different values are used for Newton's Method and Conjugate Gradient Methods, respectively $\{0.1, 0.25, 0.5, 0.9\}$ and $\{0.0001, 0.1, 0.15, 0.2\}$.

2.3.2 Golden Section Search

While golden section and fibonacci numbers are closely related, the golden section search does take much more straight forward way for calculating the new step size bounds using the approximation of $\phi = 0.618$.

Algorithm 6 Golden Section Search

1: **initialize** tolerance $l > 0$, $\alpha = 0.618$, $k = 0$, $k_{max} \in \mathbb{N}$, $(a, b) \in \mathbb{R}$.
2: $\lambda = a + (1 - \alpha)(b - a)$, $\mu = a + \alpha(b - a)$.
3: **while** $b - a > l$ **and** $k < k_{max}$ **do**
4:     **if** $\theta(\lambda) > \theta(\mu)$ **then**
5:         $a = \lambda$, $\lambda = \mu$, $\mu = a + \alpha(b - a)$
6:     **else**
7:         $b = \mu$, $\mu = \lambda$, $\lambda = a + (1 - \alpha)(b - a)$
8:     **end if**
9:     $k = k + 1$
10: **end while**
11: **return** $\lambda = \frac{a+b}{2}$

Table 3: Parameters tested for GoldenSectionSearch

Parameter	MatrixSquareSum	NegativeEntropy
a	-10, -5	-10, -5
b	5, 10	5, 10
l	1e-04, 1e-07	1e-04, 1e-07

2.3.3 Bisection Search

Bisection method is <u>commonly used line search method for differentiable functions.</u> `source`? It basically narrows down the optimal step size range using the first derivative of the target function as a direction of which bound to move next.

Algorithm 7 Bisection Search

1: **initialize** $l > 0, k = 0, k_{max} \in \mathbb{N}, (a, b) \in \mathbb{R}.$
2: **while** $|b - a| > l$ **and** $k < k_{max}$ **do**
3:     $\lambda = \frac{b+a}{2}$
4:     **if** $\theta'(\lambda) = 0$ **then**
5:        **return** λ
6:     **else if** $\theta'(\lambda) > 0$ **then**
7:        $b = \lambda$
8:     **else**
9:        $a = \lambda$
10:     **end if**
11:     $k = k + 1$
12: **end while**
13: **return** $\lambda = \frac{a+b}{2}$

Table 4: Parameters tested for BisectionSearch

Parameter	MatrixSquareSum	NegativeEntropy
a	-10, -5	-10, -5
b	5, 10	5, 10
l	1e-04, 1e-07	1e-04, 1e-07

2.3.4 Dichotomous Search

Dichotomous search is an example of sequentical searches. It uses the information of the previous evaulation of θ resulting in higher performance than e.g. uniform search.

Algorithm 8 Dichotomous Search

1: **initialize** $l > 0, k = 0, k_{max} \in \mathbb{N}, (a, b) \in \mathbb{R}$.
2: **while** $b - a > l$ **and** $k < k_{max}$ **do**
3:     $\lambda = \frac{b+a}{2} - \epsilon, \ \mu = \frac{b+a}{2} + \epsilon$
4:     **if** $\theta(\lambda) < \theta(\mu)$ **then**
5:        $b = \mu$
6:     **else**
7:        $a = \lambda$
8:     **end if**
9:     $k = k + 1$
10: **end while**
11: **return** $\lambda = \frac{a+b}{2}$

Table 5: Parameters tested for DichotomousSearch

Parameter	MatrixSquareSum	NegativeEntropy
a	-10, -5	-10, -5
b	5, 10	5, 10
ϵ	1e-07, 1e-10	1e-06, 1e-07
l	1e-05, 1e-09	1e-04, 1e-05

Initial tests show that Dichotomous Search has problems converging with some of the test values when using Newton's Method with Negative Entropy. Because of this, a slight adjustment to test parameters is done by giving a a value of -10 for NE + NM combination.

2.3.5 Fibonacci Search

Fibonacci is an interesting numerical method for providing optimal step sizes using fibonacci numbers when calculating the bounds for the step size. [better explanation of why this works?] The calculation may be somewhat heavy due to the requirement of calculating n fibonacci numbers, where n is determined by the required accuracy i.e. tolerance. However, in some real life applications these fibonacci numbers may of course be precalculated. In this thesis I will not be using a table for the fibonacci values but instead actually calculate them using a while loop. The iteration count for finding the fibonacci numbers will not be included in the performance comparison, but the time taken to find the values is included.

or what term to use for methods that do not use derivative?

is this fine, or should I also report the fibonacci iteration count and time?

Algorithm 9 Fibonacci Search

1: **initialize** $l > 0, \epsilon > 0, k = 0, k_{max} \in \mathbb{N}, (a, b) \in \mathbb{R}, n = \min_n\{\frac{b-a}{F_n} \leq l\}$.
2: $\lambda = a + \frac{F_{n-2}}{F_n}(b - a), \ \mu = a + \frac{F_{n-1}}{F_n}(b - a)$
3: **while** $k \leq n - 1$ **and** $k < k_{max}$ **do**
4:    **if** $\theta(\lambda) < \theta(\mu)$ **then**
5:       $a = \lambda, \ \lambda = \mu, \mu = a + \frac{F_{n-k-1}}{F_{n-k}}(b - a)$
6:    **else**
7:       $b = \mu, \ \lambda = \mu, \lambda = a + \frac{F_{n-k-2}}{F_{n-k}}(b - a)$
8:    **end if**
9:    $k = k + 1$
10: **end while**
11: $\lambda = \theta(\lambda)$
12: **if** $\theta(\lambda) > \theta(\mu + \epsilon)$ **then**
13:    $a = \lambda$
14: **else**
15:    $b = \mu$
16: **end if**
17: **return** $\lambda = \frac{a+b}{2}$

The F_n represents n:th Fibonacci number. To optimize algorithm performance, the values of F_n could be pre-evaluated but for our use case we are just going to generate the Fibonacci numbers using single loop algorithm of time complexity $O(n)$.

Table 6: Parameters tested for FibonacciSearch

Parameter	MatrixSquareSum	NegativeEntropy
a	-10, -5	-10, -5
b	5, 10	5, 10
ϵ	1e-08, 1e-10, 1e-12	1e-08, 1e-10, 1e-12
l	1e-06, 1e-10	1e-06, 1e-12

2.3.6 Uniform Search

Uniform search works by dividing a constrained section of a function into intervals, then choosing the interval with the lowest value and performing the same procedure again for it. To increase the precision of the method we also increase the number of intervals we divide the newly selected interval into.

More formally: break $[a, b]$ into n uniform intervals of size δ, which leads to $n + 1$ grid points $a_0 + k\delta$ with $b = a_n = a_0 + n\delta$, $k = 0, ..., n$. [2]

In addition we are going to define n as the interval count and m as the interval count multiplier.

Algorithm 10 Uniform Search

1: **initialize** $l > 0, k = 0, (k_{max}, n) \in \mathbb{N}, (a, b, m) \in \mathbb{R}$
2: $s = \frac{b-a}{n}, \ p_{min} = a$
3: **while** $s > l$ **and** $k < k_{max}$ **do**
4:    **for** $i = 0, \dots, n$ **do**
5:       $x = a + is$
6:       **if** $\theta(x) < \theta(p_{min})$ **then**
7:          $p_{min} = x$
8:       **end if**
9:       $k = k + 1$
10:    **end for**
11:    $a = p_{min} - s, \ b = p_{min} + s$
12:    $n = \lfloor nm \rfloor, \ s = \frac{b-a}{n}$
13: **end while**
14: **return** $\lambda = \frac{a+b}{2}$

Table 7: Parameters tested for UniformSearch

Parameter	MatrixSquareSum	NegativeEntropy
a	-10, -5	-10, -5
b	5, 10	5, 10
n	5, 10, 100	5, 10, 20
m	1, 1.5, 2	1, 1.5, 2
l	1e-06, 1e-08	1e-05, 1e-06

2.3.7 Newton's Search

Newton's method is another very common example of line search methods. Similarly to bisection method, Newton's method requires the target function to be differentiable up to two times as it uses the ratio of first and second derivatives to move the optimal step size forward.

Algorithm 11 Newton's Search

1: **initialize** $l > 0, k = 0$, $k_{max} \in \mathbb{N}$, $\lambda \in \mathbb{R}$
2: **while** $|\theta'(\lambda)| > l$ **and** $k < k_{max}$ **do**
3:    $\lambda = \lambda - \frac{\theta'(\lambda)}{\theta''(\lambda)}$
4:    $k = k + 1$
5: **end while**
6: **return** λ

Table 8: Parameters tested for NewtonsSearch

Parameter	MatrixSquareSum	NegativeEntropy
λ	0.5, 1, 5, 10	0.1, 0.5, 1, 5
l	1e-07, 1e-08	1e-07, 1e-08

In addition to the parameters listed in table 8, the max iteration count was added as an extra parameter for Negative Entropy function with options being either 100 or 1000 iterations.

2.3.8 Armijo's Search

Armijo's search is the only backtracking method of the ones included in this thesis.

include more in depth explanation of backtracking algorithms?

Algorithm 12 Armijo's Search

1: **initialize** $l > 0, k = 0, k_{max} \in \mathbb{N}, (\lambda, \alpha, \beta) \in \mathbb{R}$
2: $\theta_0 = \theta(0), \quad \theta_0' = \theta'(0)$
3: **while** $\theta(\lambda) > \theta_0 + \alpha\lambda\theta_0'$ **and** $k < k_{max}$ **do**
4: $\quad \lambda = \lambda\beta$
5: $\quad k = k + 1$
6: **end while**
7: **return** λ

Table 9: Parameters tested for ArmijoSearch

Parameter	MatrixSquareSum	NegativeEntropy
λ	0.9, 1, 1.1	0.9, 1, 1.1
α	0.1, 0.25, 0.5	0.1, 0.25, 0.5
β	0.5, 0.75, 0.9	0.5, 0.75, 0.9
l	1e-07	1e-07

3   Methods

Our goal is to determine how does the selection of line search method effect the overall performance of an optimization method. To answer this question with the resources introduced in the theory section, we first need to tackle the problem of choosing parameters.

3.1   Parameter Selection

Optimization method parameter selection is a whole large topic of its own, and it is not our main focus on this thesis. Therefore a very straightforward method for picking the parameters is going to be appropriate.

Using the same setup as when we are comparing the performances of different method combinations, we can also compare the effects of different parameters. Using the manually picked parameter values shown in the theory section, we can programmatically generate all the permutations for each parameter and optimization method combination.

3.2 Starting Point Selection

So we are going to compare the performances of each line search method and main method combination with their related parameter options. We are running the tests with 100 starting points of dimension $n = 50$. For practical reasons we are also going to use a limited space for the points to be randomized from.

Let a single starting point be $x = \{x_1, x_2, \ldots x_{50}\}$. For MatrixSquareSum the points are generated so that for each point applies $x_i \in [-10, 10] \; \forall \; i$. Because NegativeEntropy's domain is $\mathbb{R}_{>0}$, the same rule applies but with the bounds being $]0, 10]$ instead.

Since we are interested in overall performance, we prefer starting points that are random but evenly distributed so that no two point are the same, or too close to each other. To generate a point distribution which is random but evenly filled with points, we use the following algorithm.

Algorithm 13 Generating Even Distribution of Random Starting Points

---

1: **initialize** $(d_{min}, x_{min}, x_{max}) \in \mathbb{R}, (n, p) \in \mathbb{N}, \mathbf{x}_1 = \mathbf{random}_n(x_{min}, x_{max}), q = 1$
2: **while** $i < p$ **do**
3:     $\mathbf{y} = \mathbf{random}_n(x_{min}, x_{max})$
4:     ok = **true**
5:     **for** $j = 1 \ldots q$ **do**
6:         **if** $\|\mathbf{x}_j - \mathbf{y}\| < d_{min}$ **then**
7:             ok = **false**
8:             **break**
9:         **end if**
10:     **end for**
11:     **if** ok **then**
12:         $q = q + 1, \; i = i + 1$
13:         $\mathbf{x}_q = \mathbf{y}$
14:     **end if**
15: **end while**
16: **return** $\{x_1, x_2, ..., x_p\}$

The algorithm generates p points of dimension n. The generated points are stored in variables x_i. The function $\mathbf{random}_n(x_{min}, x_{max})$ generates a single point of dimension n filled with random values from the range $[x_{min}, x_{max}]$. On every iteration, we test to see if the newly generated point is closer than d_{min} to some existing point and if so, we discard the new point and try again.

This is by means not a perfect algorithm since its performance heavily depends on the d_{min} value used. The minimum distance selection depends on the point count and the available point space: too high values cause the algorithm to end up on infinite loop and too low values do not provide the even distribution we are looking for. Via trial and error the values presented in table 10 are picked.

Table 10: The minimum Eucleidean distances used for different target functions on distributions of 100 and 1000 points.

Point count	MSS d_{min}	NE d_{min}
100	56	28
1000	48	24

After successfully generating the starting points once, we are going to save them into a file from which they are loaded during future runs. This way we maintain the same level of randomness for each run.

3.3 Randomizing Target Function

In addition to randomizing the starting points, we are going to use randomized parameters for Matrix Square Sum target function. The values of parameters A, b and c are going to be randomized as described in section 2.1.3 but for each starting point separately. To control the randomness between runs, we give the function a random seed of the index of the current point. This way we end up with number of random functions equal to the number of starting points used. Combined with the saving of points, we can guarantee to get the same setup for each run.

3.4 Comparing Performance

There are multiple metrics that we could analyze, like algorithm success rate, main method or line search duration, iterations or function calls. The most important metric we are going to focus on is the success rate i.e. from how many of the 1000 starting points did the algorithm reach correct optimum. The solution is interpret as correct if the first 8 decimals of the found solution and correct one match. Since it is excepted that most of the algorithms find the solution correctly every time, we also need a secondary comparison metric. Here we choose to use the total duration of the algorithm, since it is often the most interesting value for real life use cases and it is comparable within different methods unlike for example total iteration

count. The other metrics will also be recorded and listed in the results section, but they will not be taken into consideration when scoring the methods.

3.4.1 Parameter Selection

4 Results

The results are going to be represented in a table with average performances of each line search method for each permutation of main methods and target functions. The performances are calculated using the best parameters found for each setup as described in the previous section. The algorithms are run with a similar setup as when calculating the best parameters, but this time with 1000 evenly randomized starting points instead of 100. Again for each starting point a unique target function configuration is generated in the case of MatrixSquareSum using the same randomizing rules as described in the previous section. The random seeds were again starting point indices i.e. integers from 0 to 1000.

After running the performance tests, we are left with the results introduced in the following eight tables. The row's of the table show the line search method used, the success rate of the algorithm, average duration, iteration count and function calls of the algorithm. The function call count includes all calls to the function and to its gradients or step size functions. In addition to these, the average duration and iterations taken by the line search algorithm alone is listed in the result tables. A table including all eight line search methods with their representative performances is generated for each main method and target function combination.

Table 11: Average performances of NewtonsMethod when minimizing MatrixSquareSum using different line search methods and 1000 randomly generated starting points

Line Search Name	s (%)	t (ms)	k	f_n	t_{LS} (ms)	k_{LS}
ConstantSearch	100.0	445.8	1.0	7.0	0.0	0.0
GoldenSectionSearch	100.0	445.0	1.0	101.0	2.6	47.0
BisectionSearch	99.6	535.9	1.0	35.0	89.0	28.0
DichotomousSearch	100.0	429.1	1.0	77.0	3.3	35.0
FibonacciSearch	100.0	453.2	1.0	119.0	4.5	55.0
UniformSearch	100.0	455.0	1.0	156.0	7.2	148.0
NewtonsSearch	100.0	443.5	1.0	8.0	0.5	0.0
ArmijoSearch	100.0	424.2	1.0	10.0	0.9	0.0

Table 12: Average performances of GradientDescentMethod when minimizing MatrixSquareSum using different line search methods and 1000 randomly generated starting points

Line Search Name	s (%)	t (ms)	k	f_n	t_{LS} (ms)	k_{LS}
ConstantSearch	98.9	4681.2	2942.0	8828.9	141.2	0.0
GoldenSectionSearch	99.4	283.7	33.1	1784.1	155.4	841.0
BisectionSearch	99.4	823.5	33.0	696.6	749.8	594.5
DichotomousSearch	99.4	270.1	33.1	1491.1	142.6	694.4
FibonacciSearch	99.4	354.3	33.1	2615.3	238.6	1223.5
UniformSearch	99.4	389.2	33.1	3508.2	282.3	3372.9
NewtonsSearch	99.4	453.3	33.1	241.2	355.7	35.3
ArmijoSearch	99.6	209.3	25.8	343.2	97.9	185.5

Table 13: Average performances of ConjugateGradientMethod when minimizing MatrixSquareSum using different line search methods and 1000 randomly generated starting points

Line Search Name	s (%)	t (ms)	k	f_n	t_{LS} (ms)	k_{LS}
ConstantSearch	99.6	989.3	300.1	1216.2	23.9	0.0
GoldenSectionSearch	100.0	429.5	50.0	2722.9	198.9	1258.5
BisectionSearch	100.0	1154.9	50.0	1106.0	1003.4	900.0
DichotomousSearch	100.0	426.3	50.0	2306.0	177.1	1050.0
FibonacciSearch	100.0	517.8	50.0	4006.0	298.5	1850.0
UniformSearch	100.0	558.2	50.0	5356.0	348.9	5100.0
NewtonsSearch	100.0	552.8	50.6	364.0	347.5	35.0
ArmijoSearch	100.0	405.3	50.0	914.8	157.6	558.8

Table 14: Average performances of HeavyBallMethod when minimizing MatrixSquareSum using different line search methods and 1000 randomly generated starting points

Line Search Name	s (%)	t (ms)	k	f_n	t_{LS} (ms)	k_{LS}
ConstantSearch	98.9	4711.5	2941.9	8828.8	134.5	0.0
GoldenSectionSearch	99.4	227.7	24.3	1303.4	127.1	613.8
BisectionSearch	99.4	666.6	24.4	514.9	607.0	438.8
DichotomousSearch	99.4	209.4	24.4	904.9	99.3	414.4
FibonacciSearch	99.4	353.2	32.7	2585.0	239.5	1209.3
UniformSearch	99.4	322.4	24.4	2587.4	232.0	2486.9
NewtonsSearch	99.4	375.1	24.4	179.6	295.0	26.4
ArmijoSearch	99.6	217.3	25.8	343.1	99.6	185.4

Table 15: Average performances of NewtonsMethod when minimizing NegativeEntropy using different line search methods and 1000 randomly generated starting points

Line Search Name	s (%)	t (ms)	k	f_n	t_{LS} (ms)	k_{LS}
ConstantSearch	100.0	3718.1	10.7	45.7	1.0	0.0
GoldenSectionSearch	100.0	2038.3	5.6	426.9	29.4	200.7
BisectionSearch	100.0	2191.6	5.6	116.9	161.4	91.5
DichotomousSearch	100.0	2231.2	6.0	265.5	52.0	119.2
FibonacciSearch	100.0	2100.7	5.6	738.5	59.0	350.9
UniformSearch	100.0	2316.0	6.1	510.3	132.5	476.7
NewtonsSearch	100.0	2266.0	5.7	192.6	348.6	53.7
ArmijoSearch	100.0	2206.9	6.3	52.6	3.1	5.5

Table 16: Average performances of GradientDescentMethod when minimizing NegativeEntropy using different line search methods and 1000 randomly generated starting points

Line Search Name	s (%)	t (ms)	k	f_n	t_{LS} (ms)	k_{LS}
ConstantSearch	100.0	13.1	16.3	51.8	0.8	0.0
GoldenSectionSearch	100.0	53.4	8.6	400.1	47.9	185.7
BisectionSearch	100.0	92.2	8.6	158.4	87.0	129.7
DichotomousSearch	100.0	54.7	8.6	289.3	48.9	130.2
FibonacciSearch	100.0	58.4	8.6	610.8	52.6	282.4
UniformSearch	100.0	121.8	10.9	889.6	114.6	842.8
NewtonsSearch	100.0	458.8	8.8	687.6	453.8	216.5
ArmijoSearch	100.0	11.6	10.2	96.5	5.6	32.0

Table 17: Average performances of ConjugateGradientMethod when minimizing NegativeEntropy using different line search methods and 1000 randomly generated starting points

Line Search Name	s (%)	t (ms)	k	f_n	t_{LS} (ms)	k_{LS}
ConstantSearch	99.9	4393.4	163.2	663.3	4248.0	0.0
GoldenSectionSearch	100.0	539.7	100.0	4681.2	451.2	2136.6
NewtonsSearch	100.0	2801.7	100.0	4147.0	2716.9	1213.0
ArmijoSearch	100.0	198.4	99.8	1119.4	112.0	412.8
BisectionSearch	100.0	759.5	100.0	1228.3	672.9	777.8
DichotomousSearch	100.0	576.4	102.7	3496.1	485.5	1538.6
FibonacciSearch	100.0	623.7	100.0	7155.1	534.0	3273.5
UniformSearch	100.0	6371.1	100.0	8208.0	6277.0	7700.0

Table 18: Average performances of HeavyBallMethod when minimizing NegativeEntropy using different line search methods and 1000 randomly generated starting points

Line Search Name	s (%)	t (ms)	k	f_n	t_{LS} (ms)	k_{LS}
ConstantSearch	100.0	8.3	14.0	45.0	0.3	0.0
GoldenSectionSearch	100.0	92.4	17.3	808.4	80.5	376.8
BisectionSearch	100.0	131.7	11.8	223.7	124.6	185.4
DichotomousSearch	100.0	68.8	11.8	486.5	60.9	224.1
FibonacciSearch	100.0	72.1	11.8	843.4	64.3	390.8
UniformSearch	100.0	134.9	12.9	1051.1	126.6	996.3
NewtonsSearch	100.0	541.0	12.8	806.1	533.7	250.6
ArmijoSearch	100.0	15.5	14.4	113.2	6.9	23.6

4.1 Analysis

Explain the min and max values of rankings

Table 19: Average duration of the algorithms using different line search methods for each main method when finding minimum for Matrix Square Sum function. Colored per column so that darker red is worse (slower) and darker green is better (faster).

Line Search Method	NM	GDM	CGM	HBM
ConstantSearch	445,80	4 681,20	989,30	4 711,50
GoldenSectionSearch	445,00	283,70	429,50	227,70
BisectionSearch	535,90	823,50	1 154,90	666,60
DichotomousSearch	429,10	270,10	426,30	209,40
FibonacciSearch	453,20	354,30	517,80	353,20
UniformSearch	455,00	389,20	558,20	322,40
NewtonsSearch	443,50	453,30	552,80	375,10
ArmijoSearch	424,20	209,30	405,30	217,30

Table 20: Average duration of the algorithms using different line search methods for each main method when finding minimum for Negative Entropy function. Colored per column so that darker red is worse (slower) and darker green is better (faster).

Line Search Method	NM	GDM	CGM	HBM
ConstantSearch	3 718,10	13,10	4 393,40	8,30
GoldenSectionSearch	2 038,30	53,40	539,70	92,40
BisectionSearch	2 191,60	92,20	759,50	131,70
DichotomousSearch	2 231,20	54,70	576,40	68,80
FibonacciSearch	2 100,70	58,40	623,70	72,10
UniformSearch	2 316,00	121,80	6 371,10	134,90
NewtonsSearch	2 266,00	458,80	2 801,70	541,00
ArmijoSearch	2 206,90	11,60	198,40	15,50

Table 21: The average of the durations for finding the minimum for both Matrix Square Sum and Negative Entropy Functions listed separately for each main method. Colored per column so that darker red is worse (slower) and darker green is better (faster).

Line Search Method	NM	GDM	CGM	HBM
ConstantSearch	2 081,95	2 347,15	2 691,35	2 359,90
GoldenSectionSearch	1 241,65	168,55	484,60	160,05
BisectionSearch	1 363,75	457,85	957,20	399,15
DichotomousSearch	1 330,15	162,40	501,35	139,10
FibonacciSearch	1 276,95	206,35	570,75	212,65
UniformSearch	1 385,50	255,50	3 464,65	228,65
NewtonsSearch	1 354,75	456,05	1 677,25	458,05
ArmijoSearch	1 315,55	110,45	301,85	116,40

Table 22: The average of the durations for finding the minimum with all the main methods listed separately for each target function. Colored per column so that darker red is worse (slower) and darker green is better (faster).

Line Search Method	MSS	NE	OVERALL
ConstantSearch	2 706,95	2 033,23	2 370,09
GoldenSectionSearch	346,48	680,95	513,71
BisectionSearch	795,23	793,75	794,49
DichotomousSearch	333,73	732,78	533,25
FibonacciSearch	419,63	713,73	566,68
UniformSearch	431,20	2 235,95	1 333,58
NewtonsSearch	456,18	1 516,88	986,53
ArmijoSearch	314,03	608,10	461,06

Table 23: The success rates of the main methods using each line search method for finding the minimum of Matrix Square Sum function. Darker red color means lower success rate.

Line Search Method	NM	GDM	CGM	HBM
ConstantSearch	100,00	98,90	99,60	98,90
GoldenSectionSearch	100,00	99,40	100,00	99,40
BisectionSearch	99,60	99,40	100,00	99,40
DichotomousSearch	100,00	99,40	100,00	99,40
FibonacciSearch	100,00	99,40	100,00	99,40
UniformSearch	100,00	99,40	100,00	99,40
NewtonsSearch	100,00	99,40	100,00	99,40
ArmijoSearch	100,00	99,60	100,00	99,60

Table 24: The success rates of the main methods using each line search method for finding the minimum of Negative Entropy function. Darker red color means lower success rate.

Line Search Method	NM	GDM	CGM	HBM
ConstantSearch	100,00	100,00	99,90	100,00
GoldenSectionSearch	100,00	100,00	100,00	100,00
BisectionSearch	100,00	100,00	100,00	100,00
DichotomousSearch	100,00	100,00	100,00	100,00
FibonacciSearch	100,00	100,00	100,00	100,00
UniformSearch	100,00	100,00	100,00	100,00
NewtonsSearch	100,00	100,00	100,00	100,00
ArmijoSearch	100,00	100,00	100,00	100,00

In addition to just comparing the raw numbers, the boxplots in figures 1 and 2 give us valuable insight of the variance of the algorithm's performance.
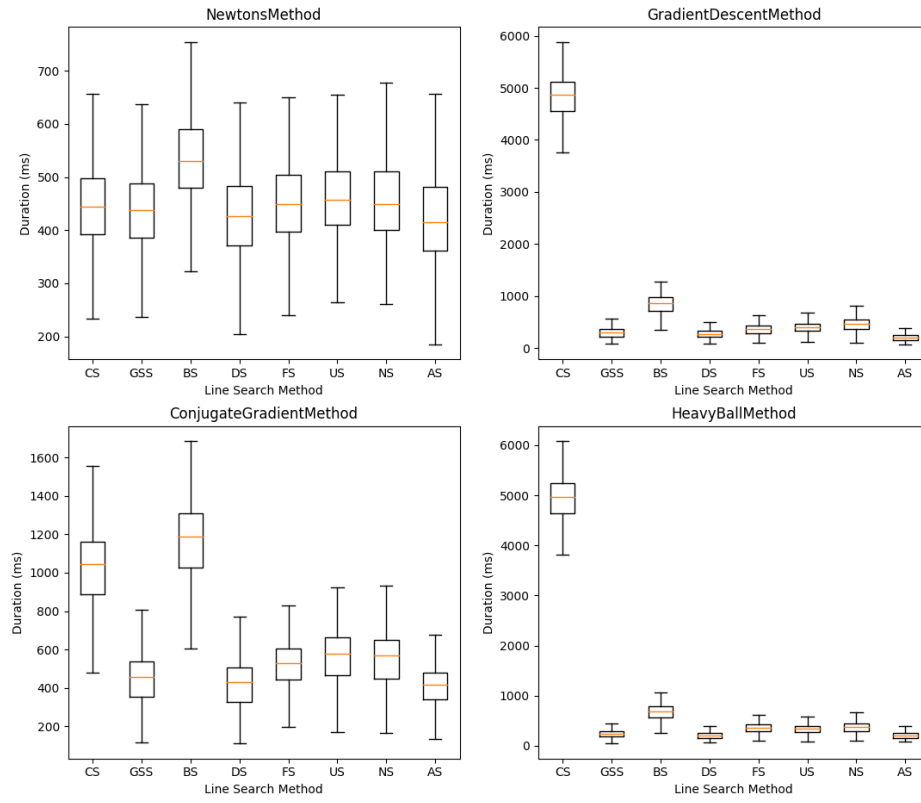
Figure 1: Box plot of solution times per line search and main method for MatrixSquareSum target function.
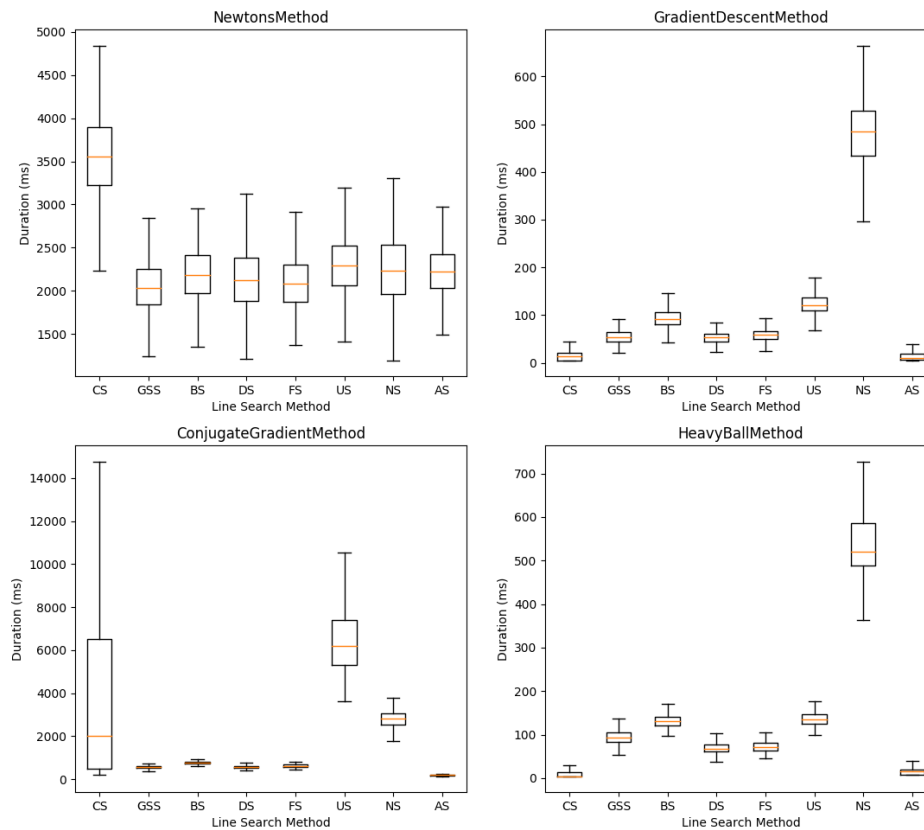
Figure 2: Box plot of solution times per line search and main method for NegativeEntropy target function.

> Explain the box plot method. Apparently the box contains 50% of the dataset and the lines rest or smthing like that?

Constant Search

Since Constant Search simply provides constant step sizes it should provide a good baseline for comparing other line search methods. The exeuction of Constant Search does basically cost zero time, so the main stress remains on the main method. This way we may be able to conclude something about the baseline performance of the main methods as well.

Looking at the overall durations of Constant Search in table 22 and the success rates at tables 23 and 24 we can quickly see that the Constant Search is the worse performing method as expected. The only scenario where constant step sizes prove themselves as rather effective are when optimizing the Negative Entropy function using Gradient Descent or Heavy Ball methods. In those cases, Constant Search was actually the best or second to best performer. This is due to the overall effectivity of gradient descent based methods for the said problem in general, and since the Constant Search is the fastest to execute, it provides a slightly better performance. This can be proven to be the case when looking at the overall duration and line search duration columns in tables 16 and 18: most of the overall duration is from the line searches.

Looking at the box plots in 1 and 2 Constant Search is clearly the worst performer by not just absolute durations but also due to the high variance of results. The only expection to this are the two previously mentioned cases when optimizing Negative Entropy with Gradient Descent based methods, where Constant Search scores one of the lowest variances.

Golden Section Search

Comparing the overall results in tables 21 and 22, we immediately notice from the solid green colored rows that Golden Section Search is one of the best performers. The success rates of this methods are also great as it does, like most methods tested, find the correct solution in almost all of the test cases.

The Golden Section Search has average variances as can be seen from box plots, which combined with one of the lowest average durations make this method one of the best of the tested line searches.

Bisection Search

Bisection Search does perform rather poorly in our test setup. It scores the second lowest total success rate, though only losing by 0.4 percentile units to others in the case of Matrix Square Sum optimized with Newton's Method. The performance is especially poor when optimizing Matrix Square Sum function and barely average with other setups, causing the overall score to be the average of the tested methods.

Judging by the box plots, the Bisection Method does also show a bit higher variances in performance than other methods especially in the case of Matrix Square Sum.

Dichotomous Search

In our first tests Dichotomous Search was one of the slower methods due to not reaching the correct minimums within the set max iterations. However, lowering the precision slightly proved to be especially helpful for this method, which ultimately turned out as one of the stronger methods, by even outperforming Golden Section Search in Matrix Square Sum optimization and scoring very similar overall scores. The success rates are also identical and reviewing variance from box plots shows no significant differences between the two methods, which share the second best place in our comparison.

Fibonacci Search

Fibonacci Search scores a slightly above average results in our performance tests, and performs much better in Negative Entropy optimization when compared to other line searche's performances. The success rates are average and variances are at the same level with Dichotomous and Golden Section Searches.

Our implementation of Fibonacci Search does compute each Fibonacci number live without storing them. This proved to not be an issue since the largest Fibonacci reached was around F_{50}. A non-official test was also performed with precomputed Fibonacci numbers, and it only improved the overall performance by approximately 10%. For real life applications precomputation should definitely be added if possible.

Uniform Search

Uniform Search is known to be slow and not very sophisticated method. This also shows in the performance as it scores averagely on Matrix Square Sum and below average for Negative Entropy. Uniform Search also struggles especially with Negative Entropy when using Conjugate Gradient Method, where it often reaches max iteration limits. This is likely due to the way we limit the domain using the Domain Limiter (algorithm 1). Overall, Uniform Search scored second lowest performances with average success rates and above average variances.

Newton's Search

Similarily to Uniform Search, Newton's Search has average performance when optimizing Matrix Square Sum but quickly falls behind in the case of Negative Entropy. Newton's Search also causes multiple warnings of max iterations reached regardless of the max iteration limit used. However, the method does eventually end up providing good enough step sizes so that success rates stayed average.

It seems that the method does not converge properly when combined with the Domain Limiter which alters the step size given by the derivatives. Because of this, a way smaller max iteration limit of just 100 or 1000 iterations proved much more effective in the case of Negative Entropy. Due to this and the use of domain limiter, the results of Newton's Search for Negative Entropy are more unreliable than in case of Matrix Square Sum. This is confirmed by the box plots as well: Newton's Search has some of the largest variances in performance of optimizing the Negative Entropy function.

Armijo's Search

Armijo's Search proves to be the overall best performing method of the tested line searches. It has the best performance in five categories out of eight, and is right on par with the top performers in the rest of the categories as well. It also has a slightly better success rate than any other method on Matrix Square Sum optimization when using Gradient Descent or Heavy Ball Method. The overall performance is also approximately 10% better than the next best overall performance of Golden Section Search. There are some differences in the variances: Armijo's Search has a slightly higher variance in Matrix Square Sum using Newton's Method but it outperforms most other method's variances when using other methods for the same function. For Negative Entropy the Armijo's Search has clearly lower variances than other methods.

4.2   Assumptions

The results introduced in this thesis contain a number of assumptions as all research usually does. The assumptions made in this paper are discussed here briefly.

One major factor to take into account when considering the results shown in this thesis is that there are only two types of different target functions that the methods are used for. Even though some variety is added by the randomization of parameters in Matrix Square Sum, both of the target functions remain relatively simple. The randomization of MSS parameters is also one factor to take into account. Randomized values are generated using different random seeds of 0 to 999 matching the index of starting points. The randomized values are selected from the range $[-0.5, 0.5]$ and the way multiplier matrix A is confirmed to be positive-definite only yields certain type of matrices. Therefore this can not be held as a totally comprehensive study of a signle type of target function either.

Another major factor to discuss is the method we use for picking values for the optimization methods. Since the problem of picking the absolute best parameters is really complex in itself, the method of picking a set of options and then comparing

the performance of each permutation separately for each setup, is justified. However, by implementing a better parameter selection, such as increasing the number of options and test cases, we can likely reach better results for each optimization methods. The tradeoff would be a higher computation time to find the optimal parameters, and while the results would highly likely improve, we can not know by how much.

5   Conclusion and Future Research

To conclude with a strict ranking for different line search methods, we should consider many more scenarios making the computation exponentially more complex. However, what we can conclude is the answer to the question introduced in the very beginning of this thesis: the choice of line search method seems to have a varying size of impact on the performance of an optimization method. Comparing the known to be worst line search methods with the best ones, there are differences of up to a single magnitude of ten. Between the more advanced methods, the differences are relatively small: half of the methods end up on a scale with just a little bit over 100 ms of difference in overall performances.

Continuing from here with a goal of generalizing the results further, I would definitely focus on increasing the number of target functions first. The only issue in this would be that the optimal parameters are likely to change with the target functions, forcing us to re-run the comparisons and find new values for all the parameters. Since this requires very high computation times, a better method for parameter selection should be implemented prior to adding new target functions. This could be something simple such as deciding a strict set of few options for each parameter but then the final results could be highly intefered by the limited parameter values.

To conclude, I would say it could be worth it to perform further research using the methods we introduce in this thesis but only if the targeted problem is something worth spending a while just to reach a potentially very slight optimization. In turn, if we would like to compare for example another set of line search methods, main methods or target functions, we should probably focus on finding better solutions to the choice of optimal parameters problems first. häiritä

References

[1] N. Andrei. An unconstrained optimization test functions collection. *Advanced Modeling and Optimization, Volume 10, Number 1, 2008*, 2000.

[2] F. Oliveira. Ms-e2122 - nonlinear optimisation. *Operations Research, Aalto University*, 2018.

6 Attachments

Constant Search

Figure 3: Best parameter permutations for different main methods using ConstantSearch for optimizing the MatrixSquareSum function with 10 sample points.

NewtonsMethod		
(λ)	s (%)	t (ms)
(1.0,)	100.0	436.3
(0.9,)	100.0	4072.7
(1.1,)	100.0	4389.6
(1.5,)	100.0	11768.2
(0.5,)	100.0	13857.0

(a) Top 5 permutations with NewtonsMethod

GradientDescentMethod		
(λ)	s (%)	t (ms)
(0.0001,)	100.0	8299.9
(0.5,)	0.0	162.8
(0.9,)	0.0	169.7
(0.1,)	0.0	187.9
(0.25,)	0.0	193.4

(b) Top 5 permutations with GradientDescentMethod

ConjugateGradientMethod		
(λ)	s (%)	t (ms)
(0.0001,)	100.0	1054.7
(0.1,)	0.0	192.4
(0.5,)	0.0	270.2
(0.25,)	0.0	280.1
(0.9,)	0.0	298.0

(c) Top 5 permutations with ConjugateGradientMethod

HeavyBallMethod		
(λ, β_{HBM})	s (%)	t (ms)
(0.0001, 0.1)	100.0	6437.8
(0.0001, 0.5)	100.0	6580.6
(0.0001, 10.0)	100.0	6714.6
(0.0001, 5.0)	100.0	6784.6
(0.0001, 1.0)	100.0	6838.7

(d) Top 5 permutations with HeavyBallMethod

Figure 4: Best parameter permutations for different main methods using ConstantSearch for optimizing the NegativeEntropy function with 10 sample points.

NewtonsMethod		
(λ)	s (%)	t (ms)
(0.9,)	100.0	3379.7
(0.5,)	100.0	7239.4
(0.25,)	100.0	17055.7
(0.1,)	100.0	46295.4

(a) Top 5 permutations with NewtonsMethod

GradientDescentMethod		
(λ)	s (%)	t (ms)
(0.5,)	100.0	17.8
(0.25,)	100.0	57.5
(0.1,)	100.0	226.9

(b) Top 5 permutations with GradientDescentMethod

ConjugateGradientMethod		
(λ)	s (%)	t (ms)
(0.1,)	100.0	2361.4
(0.0001,)	100.0	4281.1
(0.15,)	100.0	4969.4
(0.2,)	100.0	6062.7

(c) Top 5 permutations with ConjugateGradientMethod

HeavyBallMethod		
(λ, β_{HBM})	s (%)	t (ms)
(0.5, 0.1)	100.0	9.1
(0.5, 0.5)	100.0	14.7
(0.25, 0.1)	100.0	15.3
(0.25, 0.5)	100.0	15.8
(0.25, 1.0)	100.0	18.2

(d) Top 5 permutations with HeavyBallMethod

Golden Section Search

Figure 5: Best parameter permutations for different main methods using GoldenSectionSearch for optimizing the MatrixSquareSum function with 10 sample points.

NewtonsMethod		
(a, b, l)	s (%)	t (ms)
(-10, 5, 1e-07)	100.0	515.6
(-10, 10, 1e-07)	100.0	868.6
(-10, 10, 0.0001)	100.0	887.5
(-5, 5, 1e-07)	100.0	891.0
(-5, 10, 0.0001)	100.0	914.0

(a) Top 5 permutations with NewtonsMethod

GradientDescentMethod		
(a, b, l)	s (%)	t (ms)
(-10, 5, 0.0001)	100.0	279.4
(-5, 10, 0.0001)	100.0	318.0
(-5, 5, 0.0001)	100.0	344.3
(-10, 5, 1e-07)	100.0	371.1
(-10, 10, 0.0001)	100.0	379.6

(b) Top 5 permutations with GradientDescentMethod

ConjugateGradientMethod		
(a, b, l)	s (%)	t (ms)
(-5, 10, 0.0001)	100.0	436.7
(-10, 10, 0.0001)	100.0	467.8
(-5, 5, 0.0001)	100.0	489.2
(-10, 5, 0.0001)	100.0	529.8
(-10, 10, 1e-07)	100.0	533.9

(c) Top 5 permutations with ConjugateGradientMethod

HeavyBallMethod		
(a, b, l, β_{HBM})	s (%)	t (ms)
(-5, 10, 0.0001, 10.0)	100.0	230.6
(-10, 5, 0.0001, 10.0)	100.0	236.7
(-10, 10, 0.0001, 10.0)	100.0	243.6
(-5, 10, 0.0001, 5.0)	100.0	244.8
(-5, 5, 0.0001, 10.0)	100.0	252.1

(d) Top 5 permutations with HeavyBallMethod

Figure 6: Best parameter permutations for different main methods using Golden-SectionSearch for optimizing the NegativeEntropy function with 10 sample points.

NewtonsMethod		
(a, b, l)	s (%)	t (ms)
(-5, 5, 1e-07)	100.0	1758.0
(-10, 5, 0.0001)	100.0	1905.6
(-10, 10, 0.0001)	100.0	1946.7
(-10, 10, 1e-07)	100.0	1947.3
(-5, 5, 0.0001)	100.0	1968.5

(a) Top 5 permutations with NewtonsMethod

GradientDescentMethod		
(a, b, l)	s (%)	t (ms)
(-5, 5, 0.0001)	100.0	57.1
(-5, 5, 1e-07)	100.0	68.4
(-5, 10, 0.0001)	100.0	85.4
(-5, 10, 1e-07)	100.0	85.8
(-10, 5, 1e-07)	100.0	109.9

(b) Top 5 permutations with GradientDescentMethod

ConjugateGradientMethod		
(a, b, l)	s (%)	t (ms)
(-5, 5, 0.0001)	100.0	591.3
(-10, 5, 0.0001)	100.0	644.1
(-5, 5, 1e-07)	100.0	644.4
(-5, 10, 0.0001)	100.0	648.5
(-10, 5, 1e-07)	100.0	672.1

(c) Top 5 permutations with ConjugateGradientMethod

HeavyBallMethod		
(a, b, l, β_{HBM})	s (%)	t (ms)
(-5, 5, 0.0001, 0.1)	100.0	72.3
(-10, 5, 0.0001, 0.1)	100.0	75.0
(-5, 10, 0.0001, 0.1)	100.0	75.2
(-5, 5, 1e-07, 0.1)	100.0	79.7
(-5, 10, 1e-07, 0.1)	100.0	80.4

(d) Top 5 permutations with HeavyBallMethod

Bisection Search

Figure 7: Best parameter permutations for different main methods using BisectionSearch for optimizing the MatrixSquareSum function with 10 sample points.

NewtonsMethod		
(a, b, l)	s (%)	t (ms)
(-10, 5, 1e-07)	100.0	555.4
(-5, 10, 1e-07)	100.0	597.9
(-5, 5, 1e-07)	100.0	962.4
(-10, 5, 0.0001)	100.0	1025.9
(-10, 10, 0.0001)	100.0	1049.4

(a) Top 5 permutations with NewtonsMethod

GradientDescentMethod		
(a, b, l)	s (%)	t (ms)
(-10, 5, 0.0001)	100.0	836.1
(-5, 5, 0.0001)	100.0	911.2
(-10, 10, 0.0001)	100.0	921.0
(-5, 10, 0.0001)	100.0	956.1
(-5, 5, 1e-07)	100.0	1094.3

(b) Top 5 permutations with GradientDescentMethod

ConjugateGradientMethod		
(a, b, l)	s (%)	t (ms)
(-10, 10, 0.0001)	100.0	1152.1
(-5, 5, 0.0001)	100.0	1153.7
(-10, 5, 0.0001)	100.0	1157.4
(-5, 10, 0.0001)	100.0	1175.7
(-5, 10, 1e-07)	100.0	1575.7

(c) Top 5 permutations with ConjugateGradientMethod

HeavyBallMethod		
(a, b, l, β_{HBM})	s (%)	t (ms)
(-10, 10, 0.0001, 10.0)	100.0	581.2
(-5, 10, 0.0001, 5.0)	100.0	627.6
(-5, 10, 0.0001, 10.0)	100.0	676.6
(-5, 5, 0.0001, 10.0)	100.0	706.7
(-5, 5, 0.0001, 5.0)	100.0	714.3

(d) Top 5 permutations with HeavyBallMethod

Figure 8: Best parameter permutations for different main methods using Bisection Search for optimizing the NegativeEntropy function with 10 sample points.

NewtonsMethod		
(a, b, l)	s (%)	t (ms)
(-10, 10, 0.0001)	100.0	1753.5
(-10, 5, 0.0001)	100.0	2006.3
(-5, 5, 1e-07)	100.0	2044.7
(-5, 10, 0.0001)	100.0	2075.4
(-5, 5, 0.0001)	100.0	2162.0

(a) Top 5 permutations with NewtonsMethod

GradientDescentMethod		
(a, b, l)	s (%)	t (ms)
(-5, 5, 0.0001)	100.0	105.5
(-10, 5, 0.0001)	100.0	114.6
(-5, 10, 0.0001)	100.0	115.6
(-10, 5, 1e-07)	100.0	128.9
(-5, 5, 1e-07)	100.0	136.3

(b) Top 5 permutations with GradientDescentMethod

ConjugateGradientMethod		
(a, b, l)	s (%)	t (ms)
(-5, 5, 0.0001)	100.0	883.4
(-10, 10, 0.0001)	100.0	891.5
(-10, 5, 0.0001)	100.0	893.8
(-5, 10, 0.0001)	100.0	916.2
(-10, 5, 1e-07)	100.0	1092.6

(c) Top 5 permutations with ConjugateGradientMethod

HeavyBallMethod		
(a, b, l, β_{HBM})	s (%)	t (ms)
(-10, 10, 0.0001, 0.1)	100.0	138.9
(-5, 5, 0.0001, 0.1)	100.0	142.4
(-10, 5, 0.0001, 0.1)	100.0	153.4
(-5, 10, 0.0001, 0.1)	100.0	154.3
(-5, 5, 0.0001, 0.5)	100.0	182.9

(d) Top 5 permutations with HeavyBallMethod

Dichotomous Search

Figure 9: Best parameter permutations for different main methods using Dichoto-mousSearch for optimizing the MatrixSquareSum function with 10 sample points.

NewtonsMethod				GradientDescentMethod		
(a, b, ϵ, l)	s (%)	t (ms)		(a, b, ϵ, l)	s (%)	t (ms)
(-5, 10, 1e-10, 1e-09)	100.0	388.3		(-10, 5, 1e-07, 1e-05)	100.0	234.8
(-10, 10, 1e-10, 1e-09)	100.0	436.4		(-10, 5, 1e-08, 0.0001)	100.0	241.2
(-5, 5, 1e-10, 1e-09)	100.0	439.6		(-5, 5, 1e-07, 0.0001)	100.0	252.5
(-10, 5, 1e-10, 1e-09)	100.0	455.0		(-5, 10, 1e-08, 0.0001)	100.0	252.9
(-5, 10, 1e-07, 1e-05)	100.0	840.8		(-10, 5, 1e-08, 1e-05)	100.0	260.4

(a) Top 5 permutations with Newtons-Method

(b) Top 5 permutations with GradientDes-centMethod

ConjugateGradientMethod				HeavyBallMethod		
(a, b, ϵ, l)	s (%)	t (ms)		$(a, b, \epsilon, l, \beta_{HBM})$	s (%)	t (ms)
(-5, 10, 1e-07, 1e-05)	100.0	354.2		(-5, 5, 1e-07, 0.0001, 10.0)	100.0	173.6
(-5, 10, 1e-07, 0.0001)	100.0	382.0		(-10, 10, 1e-07, 0.0001, 10.0)	100.0	185.4
(-5, 5, 1e-07, 0.0001)	100.0	397.3		(-5, 10, 1e-08, 0.0001, 10.0)	100.0	193.9
(-10, 10, 1e-08, 1e-05)	100.0	406.2		(-10, 5, 1e-08, 1e-05, 5.0)	100.0	202.6
(-10, 5, 1e-08, 0.0001)	100.0	406.7		(-5, 5, 1e-07, 1e-05, 10.0)	100.0	205.2

(c) Top 5 permutations with ConjugateGra-dientMethod

(d) Top 5 permutations with HeavyBall-Method

Figure 10: Best parameter permutations for different main methods using Dichoto-mousSearch for optimizing the NegativeEntropy function with 10 sample points.

NewtonsMethod			GradientDescentMethod		
(a, b, ϵ, l)	s (%)	t (ms)	(a, b, ϵ, l)	s (%)	t (ms)
(-10, 10, 1e-07, 1e-05)	100.0	1756.0	(-5, 5, 1e-07, 0.0001)	100.0	54.3
(-10, 5, 1e-06, 1e-05)	100.0	1846.8	(-5, 10, 1e-06, 0.0001)	100.0	61.7
(-10, 10, 1e-06, 0.0001)	100.0	1882.9	(-5, 10, 1e-06, 1e-05)	100.0	63.8
(-10, 5, 1e-06, 0.0001)	100.0	1893.5	(-5, 5, 1e-07, 1e-05)	100.0	64.6
(-10, 10, 1e-06, 1e-05)	100.0	1923.0	(-5, 10, 1e-07, 0.0001)	100.0	65.4

(a) Top 5 permutations with Newtons-Method

(b) Top 5 permutations with GradientDes-centMethod

ConjugateGradientMethod			HeavyBallMethod		
(a, b, ϵ, l)	s (%)	t (ms)	$(a, b, \epsilon, l, \beta_{HBM})$	s (%)	t (ms)
(-5, 5, 1e-06, 0.0001)	100.0	640.6	(-5, 5, 1e-06, 1e-05, 0.1)	100.0	70.4
(-10, 5, 1e-06, 0.0001)	100.0	642.7	(-5, 5, 1e-06, 0.0001, 0.1)	100.0	70.8
(-5, 5, 1e-07, 0.0001)	100.0	645.3	(-10, 5, 1e-07, 0.0001, 0.1)	100.0	73.3
(-5, 10, 1e-07, 0.0001)	100.0	657.0	(-5, 10, 1e-06, 0.0001, 0.1)	100.0	74.0
(-5, 5, 1e-06, 1e-05)	100.0	680.9	(-5, 5, 1e-07, 0.0001, 0.1)	100.0	74.0

(c) Top 5 permutations with ConjugateGra-dientMethod

(d) Top 5 permutations with HeavyBall-Method

Fibonacci Search

Figure 11: Best parameter permutations for different main methods using FibonacciSearch for optimizing the MatrixSquareSum function with 10 sample points.

NewtonsMethod		
(a, b, ϵ, l)	s (%)	t (ms)
(-5, 5, 1e-10, 1e-10)	100.0	443.2
(-5, 5, 1e-12, 1e-10)	100.0	452.3
(-5, 5, 1e-08, 1e-10)	100.0	459.6
(-10, 5, 1e-10, 1e-10)	100.0	460.9
(-10, 5, 1e-12, 1e-10)	100.0	469.7

(a) Top 5 permutations with NewtonsMethod

GradientDescentMethod		
(a, b, ϵ, l)	s (%)	t (ms)
(-5, 10, 1e-12, 1e-06)	100.0	355.6
(-5, 10, 1e-10, 1e-06)	100.0	356.0
(-5, 5, 1e-12, 1e-06)	100.0	366.8
(-10, 10, 1e-12, 1e-06)	100.0	369.2
(-10, 5, 1e-12, 1e-06)	100.0	369.4

(b) Top 5 permutations with GradientDescentMethod

ConjugateGradientMethod		
(a, b, ϵ, l)	s (%)	t (ms)
(-5, 10, 1e-10, 1e-06)	100.0	473.8
(-10, 10, 1e-10, 1e-06)	100.0	485.0
(-5, 10, 1e-12, 1e-06)	100.0	490.3
(-10, 5, 1e-08, 1e-06)	100.0	508.6
(-10, 10, 1e-12, 1e-06)	100.0	516.0

(c) Top 5 permutations with ConjugateGradientMethod

HeavyBallMethod		
$(a, b, \epsilon, l, \beta_{HBM})$	s (%)	t (ms)
(-10, 5, 1e-12, 1e-06, 10.0)	100.0	251.4
(-5, 5, 1e-12, 1e-06, 10.0)	100.0	268.5
(-5, 5, 1e-08, 1e-06, 10.0)	100.0	273.0
(-10, 10, 1e-12, 1e-06, 10.0)	100.0	273.5
(-5, 5, 1e-10, 1e-06, 10.0)	100.0	273.5

(d) Top 5 permutations with HeavyBallMethod

Figure 12: Best parameter permutations for different main methods using FibonacciSearch for optimizing the NegativeEntropy function with 10 sample points.

NewtonsMethod		
(a, b, ϵ, l)	s (%)	t (ms)
(-5, 5, 1e-12, 1e-12)	100.0	1834.0
(-10, 5, 1e-08, 1e-06)	100.0	1908.9
(-10, 10, 1e-12, 1e-06)	100.0	1946.8
(-5, 10, 1e-10, 1e-06)	100.0	1957.9
(-10, 10, 1e-10, 1e-06)	100.0	1967.3

(a) Top 5 permutations with NewtonsMethod

GradientDescentMethod		
(a, b, ϵ, l)	s (%)	t (ms)
(-5, 5, 1e-12, 1e-06)	100.0	74.5
(-5, 10, 1e-08, 1e-06)	100.0	75.8
(-5, 5, 1e-10, 1e-06)	100.0	78.0
(-5, 5, 1e-08, 1e-06)	100.0	78.3
(-10, 5, 1e-12, 1e-12)	100.0	78.6

(b) Top 5 permutations with GradientDescentMethod

ConjugateGradientMethod		
(a, b, ϵ, l)	s (%)	t (ms)
(-5, 5, 1e-08, 1e-06)	100.0	682.8
(-5, 5, 1e-12, 1e-06)	100.0	683.5
(-5, 5, 1e-10, 1e-06)	100.0	701.5
(-10, 5, 1e-10, 1e-06)	100.0	728.0
(-5, 10, 1e-08, 1e-06)	100.0	729.5

(c) Top 5 permutations with ConjugateGradientMethod

HeavyBallMethod		
$(a, b, \epsilon, l, \beta_{HBM})$	s (%)	t (ms)
(-5, 5, 1e-12, 1e-06, 0.1)	100.0	76.4
(-5, 5, 1e-10, 1e-06, 0.1)	100.0	82.2
(-10, 5, 1e-08, 1e-06, 0.1)	100.0	85.1
(-5, 10, 1e-10, 1e-06, 0.1)	100.0	86.9
(-5, 5, 1e-08, 1e-06, 0.1)	100.0	87.6

(d) Top 5 permutations with HeavyBallMethod

Uniform Search

Figure 13: Best parameter permutations for different main methods using Uniform-Search for optimizing the MatrixSquareSum function with 10 sample points.

NewtonsMethod		
(a, b, n, m, l)	s (%)	t (ms)
(-10, 5, 5, 1.5, 1e-06)	100.0	424.0
(-5, 5, 10, 1.5, 1e-06)	100.0	424.5
(-5, 5, 10, 1, 1e-08)	100.0	449.1
(-10, 10, 5, 1.5, 1e-08)	100.0	459.5
(-5, 10, 5, 1, 1e-06)	100.0	463.1

(a) Top 5 permutations with NewtonsMethod

GradientDescentMethod		
(a, b, n, m, l)	s (%)	t (ms)
(-10, 5, 5, 1, 1e-06)	100.0	231.0
(-5, 5, 5, 1, 1e-06)	100.0	254.0
(-5, 10, 5, 1, 1e-06)	100.0	287.5
(-5, 10, 10, 1, 1e-06)	100.0	289.1
(-5, 5, 5, 1, 1e-08)	100.0	381.7

(b) Top 5 permutations with GradientDescentMethod

ConjugateGradientMethod		
(a, b, n, m, l)	s (%)	t (ms)
(-5, 10, 5, 1, 1e-06)	100.0	328.1
(-10, 10, 5, 1, 1e-06)	100.0	364.8
(-5, 5, 5, 1, 1e-06)	100.0	371.5
(-10, 5, 5, 1, 1e-06)	100.0	378.8
(-5, 5, 10, 1, 1e-06)	100.0	529.3

(c) Top 5 permutations with ConjugateGradientMethod

HeavyBallMethod		
$(a, b, n, m, l, \beta_{HBM})$	s (%)	t (ms)
(-10, 10, 5, 1, 1e-06, 10.0)	100.0	137.3
(-5, 10, 5, 1, 1e-06, 10.0)	100.0	155.6
(-5, 5, 5, 1, 1e-06, 10.0)	100.0	170.2
(-10, 10, 5, 1, 1e-06, 5.0)	100.0	185.4
(-5, 5, 5, 1, 1e-06, 5.0)	100.0	200.0

(d) Top 5 permutations with HeavyBallMethod

Figure 14: Best parameter permutations for different main methods using Uniform-Search for optimizing the NegativeEntropy function with 10 sample points.

NewtonsMethod		
(a, b, n, m, l)	s (%)	t (ms)
(0, 5, 5, 1, 1e-05)	100.0	1339.2
(-5, 5, 5, 1, 1e-05)	100.0	2047.3
(-5, 10, 5, 1, 1e-05)	100.0	2059.8
(0, 5, 5, 1, 1e-06)	100.0	2068.5
(0, 5, 10, 1, 1e-06)	100.0	2125.3

(a) Top 5 permutations with NewtonsMethod

GradientDescentMethod		
(a, b, n, m, l)	s (%)	t (ms)
(0, 5, 10, 1, 1e-05)	100.0	139.2
(0, 5, 10, 1.5, 1e-06)	100.0	159.3
(0, 5, 10, 1.5, 1e-05)	100.0	160.6
(0, 5, 10, 1, 1e-06)	100.0	161.2
(0, 5, 10, 2, 1e-06)	100.0	169.0

(b) Top 5 permutations with GradientDescentMethod

ConjugateGradientMethod		
(a, b, n, m, l)	s (%)	t (ms)
(0, 5, 10, 1, 1e-05)	100.0	6718.5
(0, 5, 5, 1, 1e-05)	100.0	7035.3
(-5, 5, 10, 1, 1e-05)	100.0	7444.3
(-5, 5, 5, 1, 1e-05)	100.0	7807.4
(0, 5, 5, 1.5, 1e-05)	100.0	8299.8

(c) Top 5 permutations with ConjugateGradientMethod

HeavyBallMethod		
$(a, b, n, m, l, \beta_{HBM})$	s (%)	t (ms)
(0, 5, 10, 1, 1e-05, 0.1)	100.0	138.5
(0, 5, 5, 1.5, 1e-05, 0.1)	100.0	154.5
(0, 5, 5, 1, 1e-05, 0.1)	100.0	166.8
(0, 5, 10, 1, 1e-06, 0.1)	100.0	177.9
(0, 5, 10, 1.5, 1e-05, 0.1)	100.0	178.0

(d) Top 5 permutations with HeavyBallMethod

Newton's Search

Figure 15: Best parameter permutations for different main methods using NewtonsSearch for optimizing the MatrixSquareSum function with 10 sample points.

NewtonsMethod		
(λ, l)	s (%)	t (ms)
(1, 1e-07)	100.0	411.3
(10, 1e-08)	100.0	435.8
(1, 1e-08)	100.0	445.1
(5, 1e-08)	100.0	447.9
(5, 1e-07)	100.0	449.2

(a) Top 5 permutations with NewtonsMethod

GradientDescentMethod		
(λ, l)	s (%)	t (ms)
(1, 1e-08)	100.0	412.0
(5, 1e-07)	100.0	415.6
(10, 1e-07)	100.0	430.2
(5, 1e-08)	100.0	460.3
(10, 1e-08)	100.0	470.8

(b) Top 5 permutations with GradientDescentMethod

ConjugateGradientMethod		
(λ, l)	s (%)	t (ms)
(1, 1e-07)	100.0	502.4
(0.5, 1e-07)	100.0	537.2
(10, 1e-08)	100.0	554.9
(5, 1e-07)	100.0	569.9
(1, 1e-08)	100.0	592.0

(c) Top 5 permutations with ConjugateGradientMethod

HeavyBallMethod		
$(\lambda, l, \beta_{HBM})$	s (%)	t (ms)
(10, 1e-07, 10.0)	100.0	338.8
(5, 1e-07, 10.0)	100.0	360.2
(10, 1e-08, 5.0)	100.0	361.5
(0.5, 1e-07, 10.0)	100.0	361.9
(1, 1e-08, 10.0)	100.0	382.9

(d) Top 5 permutations with HeavyBallMethod

Figure 16: Best parameter permutations for different main methods using NewtonsSearch for optimizing the NegativeEntropy function with 10 sample points.

NewtonsMethod		
(λ, l)	s (%)	t (ms)
(0.5, 1e-08, 100)	100.0	1659.8
(1, 1e-08, 100)	100.0	1662.6
(1, 1e-07, 100)	100.0	1671.8
(5, 1e-07, 100)	100.0	1707.0
(0.5, 1e-07, 100)	100.0	1723.9

(a) Top 5 permutations with NewtonsMethod

GradientDescentMethod		
(λ, l)	s (%)	t (ms)
(0.5, 1e-07, 100)	100.0	476.2
(0.5, 1e-08, 100)	100.0	501.1
(1, 1e-08, 100)	100.0	501.8
(0.1, 1e-08, 100)	100.0	520.6
(5, 1e-07, 100)	100.0	521.0

(b) Top 5 permutations with GradientDescentMethod

ConjugateGradientMethod		
(λ, l)	s (%)	t (ms)
(0.1, 1e-08, 100)	100.0	2779.5
(0.1, 1e-07, 100)	100.0	2789.7
(0.5, 1e-07, 100)	100.0	2921.1
(0.5, 1e-08, 100)	100.0	2985.0
(1, 1e-07, 100)	100.0	3050.7

(c) Top 5 permutations with ConjugateGradientMethod

HeavyBallMethod		
$(\lambda, l, \beta_{HBM})$	s (%)	t (ms)
(1, 1e-07, 100, 0.1)	100.0	586.3
(0.1, 1e-08, 100, 0.1)	100.0	591.6
(1, 1e-08, 100, 0.1)	100.0	602.4
(0.5, 1e-07, 100, 0.1)	100.0	609.2
(0.1, 1e-07, 100, 0.1)	100.0	614.9

(d) Top 5 permutations with HeavyBallMethod

Armijo's Search

Figure 17: Best parameter permutations for different main methods using ArmijoSearch for optimizing the MatrixSquareSum function with 10 sample points.

NewtonsMethod		
(λ, α, β)	s (%)	t (ms)
(1, 0.1, 0.5)	100.0	429.7
(1, 0.25, 0.75)	100.0	484.5
(1, 0.25, 0.5)	100.0	493.8
(1, 0.25, 0.9)	100.0	493.9
(1, 0.1, 0.9)	100.0	498.7

(a) Top 5 permutations with NewtonsMethod

GradientDescentMethod		
(λ, α, β)	s (%)	t (ms)
(1.1, 0.25, 0.5)	100.0	143.0
(1.1, 0.5, 0.5)	100.0	211.3
(1, 0.25, 0.5)	100.0	215.6
(0.9, 0.25, 0.5)	100.0	216.4
(0.9, 0.5, 0.5)	100.0	236.4

(b) Top 5 permutations with GradientDescentMethod

ConjugateGradientMethod		
(λ, α, β)	s (%)	t (ms)
(1.1, 0.25, 0.5)	100.0	262.6
(1, 0.25, 0.5)	100.0	339.0
(1.1, 0.1, 0.5)	100.0	353.9
(1, 0.25, 0.75)	100.0	361.5
(0.9, 0.25, 0.5)	100.0	376.6

(c) Top 5 permutations with ConjugateGradientMethod

HeavyBallMethod		
$(\lambda, \alpha, \beta, \beta_{HBM})$	s (%)	t (ms)
(1.1, 0.25, 0.5, 0.1)	100.0	142.5
(0.9, 0.25, 0.5, 10.0)	100.0	151.8
(1, 0.25, 0.5, 10.0)	100.0	165.7
(1, 0.25, 0.5, 0.5)	100.0	170.4
(1.1, 0.5, 0.75, 10.0)	100.0	173.8

(d) Top 5 permutations with HeavyBallMethod

Figure 18: Best parameter permutations for different main methods using ArmijoSearch for optimizing the NegativeEntropy function with 10 sample points.

NewtonsMethod		
(λ, α, β)	s (%)	t (ms)
(1.1, 0.5, 0.9)	100.0	2094.1
(1, 0.5, 0.75)	100.0	2457.5
(1, 0.5, 0.5)	100.0	2659.1
(1, 0.5, 0.9)	100.0	2756.4
(1.1, 0.25, 0.75)	100.0	2854.2

(a) Top 5 permutations with NewtonsMethod

GradientDescentMethod		
(λ, α, β)	s (%)	t (ms)
(1.1, 0.5, 0.75)	100.0	11.2
(1.1, 0.5, 0.9)	100.0	12.6
(0.9, 0.5, 0.9)	100.0	12.7
(1.1, 0.25, 0.75)	100.0	13.9
(1.1, 0.25, 0.5)	100.0	15.4

(b) Top 5 permutations with GradientDescentMethod

ConjugateGradientMethod		
(λ, α, β)	s (%)	t (ms)
(0.9, 0.5, 0.5)	100.0	221.7
(1.1, 0.5, 0.5)	100.0	242.8
(0.9, 0.5, 0.75)	100.0	252.2
(1.1, 0.5, 0.75)	100.0	255.0
(1, 0.5, 0.5)	100.0	262.5

(c) Top 5 permutations with ConjugateGradientMethod

HeavyBallMethod		
$(\lambda, \alpha, \beta, \beta_{HBM})$	s (%)	t (ms)
(1, 0.5, 0.5, 0.1)	100.0	11.4
(1.1, 0.25, 0.5, 0.1)	100.0	14.1
(0.9, 0.5, 0.75, 0.1)	100.0	14.5
(0.9, 0.25, 0.5, 0.1)	100.0	14.5
(1, 0.5, 0.9, 0.1)	100.0	15.8

(d) Top 5 permutations with HeavyBallMethod